

MPI Cluster Setup – Multi-node – MPICH2

*** PLEASE **DO NOT TRY THIS** IN LAB 6019. THE CLUSTER IS ALREADY SET UP ***

The file system used for the cluster setup is network file system (NFS). There is a conceptual master storage location, that is mounted over a location in each local node, for providing a view of globally shared data storage.

Here a new user account is created at each node with name **mpiuser**, creation of which is discussed later (not necessary, you may use your existing account).

1. Defining hostnames

Edit /etc/hosts file to provide names for nodes to be connected for ease of use, in all the nodes.

An example /etc/hosts file may look as follows

```
127.0.0.1      localhost
192.168.133.100  mpimaster
192.168.133.101  mpi1
192.168.133.102  mpi2
192.168.133.103  mpi3
```

2. Installing NFS

NFS allows to create a folder at notional master and sync it over other nodes

At master node

```
sudo apt-get install nfs-server
```

At other nodes

```
sudo apt-get install nfs-client
```

3. Sharing Master Folder

Make a folder in all nodes, where we will store all our data and programs

```
sudo mkdir /mirror
```

Modify the /etc/exports file in Master node to allow sharing of contents in Master's mirror folder to other nodes

Way – 1 : From terminal

```
echo "/mirror *(rw, sync)" | sudo tee -a /etc/exports
```

Way – 2 : Open /etc/exports and add the following line at the end

```
/mirror *(rw, sync)
```

Now at Master node, restart the NFS kernel server

```
sudo service nfs-kernel-server restart
```

4. Mounting mirror folder in nodes

Add the following line in /etc/fstab file in non-master nodes
mpimaster:/mirror /mirror nfs
assuming hostname of master node given as mpimaster

5. Define a user for running MPI programs

Define a new user, say mpiuser, on all nodes with same name and same userid with home directory in /mirror

```
sudo useradd -m -d /mirror/mpiuser -s /bin/bash -uid 1001 -G sudo mpiuser
```

Setup password

```
sudo passwd mpiuser
```

Also give the ownership of /mirror to mpiuser

```
sudo chown mpiuser /mirror
```

6. Install SSH server

Install OpenSSH server in all nodes

```
sudo apt-get install openssh-server
```

7. Setting up passwordless SSH for communication between nodes

At master node, login to mpiuser

```
su -mpiuser
```

Generate RSA key pair

```
ssh-keygen -t rsa
```

Add this key to authorized keys

```
cd .ssh
```

```
cat id_rsa.pub >> authorized_keys
```

You may check now if passwordless ssh is working by

```
ssh non-master_node_hostname
```

where non-master_node_hostname is the hostname of any of the other node mentioned in /etc/hosts file

8. Install MPICH2

Install MPICH2 on all the nodes

```
sudo apt-get install mpich2
```

Alternatively, you may install openmpi in one node in the cluster anywhere in the shared /mirror directory

Test if successful by

```
which mpiexec
```

```
which mpirun
```

9. Testing

Write a simple MPI program, say pgm1.c and compile as

```
mpicc pgm1.c -o pgm1
```

Run the program in multiple nodes as

```
mpirun -np 10 -host mpimaster,mpi1,mpi2 ./pgm1
```

where -np 10 says to spawns 10 processes in 3 nodes

Run program in single node as

```
Mpirun -np 4 ./pgm1
```

where -np 4 says to spawn 4 processes in the local node

Reference : <https://help.ubuntu.com/community/MpichCluster>